


ICS 321 Fall 2012

SQL in a Server Environment (ii)

Asst. Prof. Lipyeow Lim
Information & Computer Science Department
University of Hawaii at Manoa

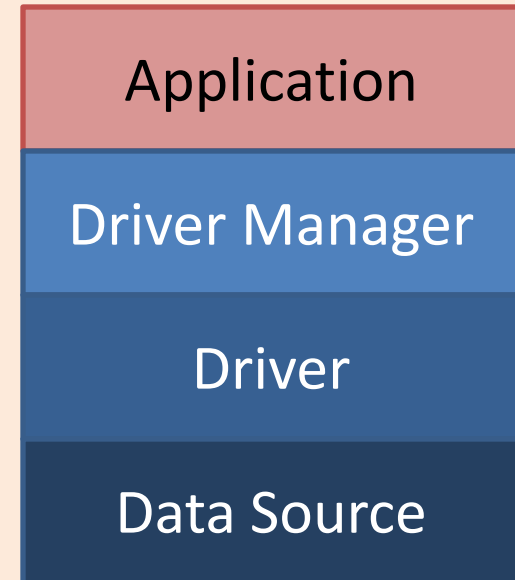
Alternative to Embedded SQL

- What if we want to compile an application without the need for a DBMS-specific pre-compiler ?
- Use a library of database calls
 - Standardized (non-DBMS-specific) API
 - Pass SQL-strings from host language and presents result sets in a language friendly way
 - Eg. ODBC for C/C++ and JDBC for Java
 - DBMS-neutral
 - A driver traps the calls and translates them into DBMS-specific code



ODBC/JDBC Architecture

- Application
 - Initiates connections
 - Submits SQL statements
 - Terminates connections
- Driver Manager
 - Loads the right JDBC driver
- Driver
 - Connects to the data source,
 - Transmit requests,
 - Returns results and error codes
- Data Source
 - DBMS



4 Types of Drivers

- Type I: Bridge
 - Translate SQL commands to non-native API
 - eg. JDBC-ODBC bridge. JDBC is translated to ODBC to access an ODBC compliant data source.
- Type II: Direct Translation to native API via non-Java driver
 - Translates SQL to native API of data source.
 - Needs DBMS-specific library on each client.
- Type III: Network bridge
 - SQL stmts sent to a middleware server that talks to the data source. Hence small JDBC driver at each client
- Type IV: Direct Translation to native API via Java driver
 - Converts JDBC calls to network protocol used by DBMS.
 - Needs DBMS-specific Java driver at each client.

High Level Steps

1. Load the ODBC/JDBC driver
2. Connect to the data source
3. [optional] Prepare the SQL statements
4. Execute the SQL statements
5. Iterate over the resultset
6. Close the connection

Getting Data to/fro Host Language

- No declaration of shared variables
- Variables in host language is bound to columns of a SQL cursor
- ODBC
 - SQLBindCol – gets data from SQL environment to host variables.
 - SQLBindParameter – gets data from host variables to SQL environment
- JDBC
 - ResultSet class
 - PreparedStatement class

Prepare Statement or Not ?

```
String sql="SELECT * FROM books WHERE price < ?";  
PreparedStatement pstmt = conn.prepareStatement(sql);  
Pstmt.setFloat(1, usermaxprice);  
Pstmt.executeUpdate();
```

- Executing without preparing statement
 - After DBMS receives SQL statement,
 - The SQL is compiled,
 - An execution plan is chosen by the optimizer,
 - The execution plan is evaluated by the DBMS engine
 - The results are returned
- `conn.prepareStatement`
 - Compiles and picks an execution plan
- `pstmt.executeUpdate`
 - Evaluates the execution plan with the parameters and gets the results

cf. Static vs
Dynamic
SQL

ResultSet

```
ResultSet rs = stmt.executeQuery(sqlstr);
while( rs.next() ){
    col1val = rs.getString(1); ...
}
```

- Iterate over the results of a SQL statement -- cf. cursor
- Note that types of column values do not need to be known at compile time

| SQL Type | Java Class | accessor |
|------------------|--------------------|--------------|
| BIT | Boolean | getBoolean |
| CHAR, VARCHAR | String | getString |
| DOUBLE, FLOAT | Double | getDouble |
| INTEGER | Integer | getInt |
| REAL | Double | getFloat |
| DATE | Java.sql.Date | getDate |
| TIME | Java.sql.Time | getTime |
| TIMESTAMP | Java.sql.TimeStamp | getTimestamp |

RowSet

- When inserting lots of data, calling an execute statement for each row can be inefficient
 - A message is sent for each execute
- Many APIs provide a rowset implementation
 - A set of rows is maintained in-memory on the client
 - A single execute will then insert the set of rows in a single message
- Pros: high performance
- Cons: data can be lost if client crashes.
- Analogous rowset for reads (ie. ResultSet) also available

Stored Procedures

- What ?
 - A procedure that is called and executed via a single SQL statement
 - Executed in the same process space of the DBMS server
 - Can be programmed in SQL, C, java etc
 - The procedure is stored within the DBMS
- Advantages:
 - Encapsulate application logic while staying close to the data
 - Re-use of application logic by different users
 - Avoid tuple-at-a-time return of records through cursors

SQL Stored Procedures

CREATE PROCEDURE ShowNumReservations

SELECT S.sid, S.sname, COUNT(*)

FROM Sailors S, Reserves R

WHERE S.sid = R.sid

GROUP BY S.sid, S.sname

- Parameters modes: IN, OUT, INOUT

CREATE PROCEDURE IncreaseRating (IN sailor_sid
INTEGER, IN increase INTEGER)

UPDATE Sailors

SET rating = rating + increase

WHERE sid = sailor_sid

Java Stored Procedures

```
CREATE PROCEDURE TopSailors (  
    IN num INTEGER)  
LANGUAGE JAVA  
EXTERNAL NAME  
    “file:///c:/storedProcs/rank.jar”
```

Calling Stored Procedures

- SQL: **CALL** IncreaseRating(101, 2);
- Embedded SQL in C:
EXEC SQL BEGIN DECLARE SECTION
int sid; int rating;
EXEC SQL END DECLARE SECTION
EXEC SQL CALL IncreaseRating(:sid, :rating);
- JDBC
CallableStatement cstmt = conn.prepareCall("{call Show Sailors}");
ResultSet rs=cstmt.executeQuery();
- ODBC
SQLCHAR *stmt = (SQLCHAR *)"CALL ShowSailors";
cliRC = SQLPrepare(hstmt, stmt, SQL_NTS);
cliRC = SQLExecute(hstmt);

User Defined Functions (UDFs)

- Extend and add to the support provided by SQL built-in functions
- Three types of UDFs
 - **Scalar**: returns a single-valued answer. Eg. Built-in SUBSTR()
 - **Column**: returns a single-valued answer from a column of values. Eg. AVG()
 - **Table**: returns a table. Invoked in the FROM clause.
- Programmable in SQL, C, JAVA.

Scalar UDFs

- Returns the tangent of a value

```
CREATE FUNCTION TAN (X DOUBLE)  
RETURNS DOUBLE  
LANGUAGE SQL  
CONTAINS SQL  
RETURN SIN(X)/COS(X)
```

- Reverses a string

```
CREATE FUNCTION REVERSE(INSTR  
  VARCHAR(4000))  
RETURNS VARCHAR(4000)  
CONTAINS SQL
```

BEGIN ATOMIC

```
DECLARE REVSTR, RESTSTR  
  VARCHAR(4000) DEFAULT '';  
DECLARE LEN INT;  
IF INSTR IS NULL THEN  
  RETURN NULL;  
END IF;  
SET (RESTSTR, LEN) = (INSTR,  
  LENGTH(INSTR));  
WHILE LEN > 0 DO  
  SET (REVSTR, RESTSTR, LEN)  
    = (SUBSTR(RESTSTR, 1, 1) CONCAT  
      REVSTR, SUBSTR(RESTSTR, 2, LEN  
        - 1), LEN - 1);  
END WHILE;  
RETURN REVSTR;  
END
```

Table UDFs

- returns the employees in a specified department number.

```
CREATE FUNCTION DEPTEMPLOYEES (DEPTNO CHAR(3))
```

```
RETURNS TABLE (
```

```
    EMPNO CHAR(6),
```

```
    LASTNAME VARCHAR(15),
```

```
    FIRSTNAME VARCHAR(12))
```

```
LANGUAGE SQL
```

```
READS SQL DATA
```

```
RETURN
```

```
    SELECT EMPNO, LASTNAME, FIRSTNME
```

```
    FROM EMPLOYEE
```

```
    WHERE EMPLOYEE.WORKDEPT
```

```
        = DEPTEMPLOYEES.DEPTNO
```


Java UDFs

```
CREATE FUNCTION tableUDF ( DOUBLE ) RETURNS TABLE (
    name VARCHAR(20),
    job VARCHAR(20),
    salary DOUBLE )
EXTERNAL NAME 'MYJAR1:UDFsrv!tableUDF'
LANGUAGE JAVA
PARAMETER STYLE DB2GENERAL
NOT DETERMINISTIC
FENCED
NO SQL
NO EXTERNAL ACTION
SCRATCHPAD 10
FINAL CALL
DISALLOW PARALLEL
NO DBINFO@
```

import COM.ibm.db2.app.UDF;

```
public void tableUDF(
    double inSalaryFactor,
    String outName,
    String outJob,
    double outNewSalary)
    throws Exception
{
    int intRow = 0;
    ...
} // tableUDF } // UDFsrv class
```