

ICS 321 Spring 2013


# Overview of Storage & Indexing (ii)

Asst. Prof. Lipyeow Lim  
Information & Computer Science Department  
University of Hawaii at Manoa

# Indexes


- An index on a file speeds up selections on the *search key fields* for the index.
  - Any subset of the fields of a relation can be the search key for an index on the relation.
  - *Search key* is *not* the same as *key* (minimal set of fields that uniquely identify a record in a relation).
- An index contains a collection of *data entries*, and supports efficient retrieval of all data entries  $k^*$  with a given key value  $k$ .
  - A data entry is usually in the form  $\langle \text{key}, \text{rid} \rangle$
  - Given data entry  $k^*$ , we can find record with key  $k$  in at most one disk I/O. (Details soon ...)

# B+ Tree Indexes



- Leaf pages contain **data entries**, and are chained (prev & next)
- A data entry typically contain a key value and a rid.
- Non-leaf pages have **index entries**; only used to direct searches:

# Example B+ Tree




- Find 28\*? 29\*? All  $> 15^*$  and  $< 30^*$
- Insert/delete: Find data entry in leaf, then change it. Need to adjust parent sometimes.
  - And change sometimes bubbles up the tree

# Point Queries using B+ Trees

SELECT \*  
FROM Employees  
WHERE age=30

Assume heap file  
data storage

- Use index to find 30\*
- Request tuple from buffer manager
- If not in bufferpool, fetch page from disk




# Range Queries using B+ Trees


SELECT \*  
FROM Employees  
WHERE age > 30

Assume heap file  
data storage

- Use index to find 30\*
- For each data entry to the right of 30\*
- Request tuples from buffer manager
- If not in bufferpool, fetch page from disk



# Hash-Based Indexes



- Index is a collection of *buckets* that contain data entries
  - Bucket = *primary page* plus zero or more *overflow pages*.
- *Hashing function h*:  $h(r)$  = bucket in which (data entry for) record  $r$  belongs.  $h$  looks at the *search key* fields of  $r$ .
- *No “index entries” in this scheme.*

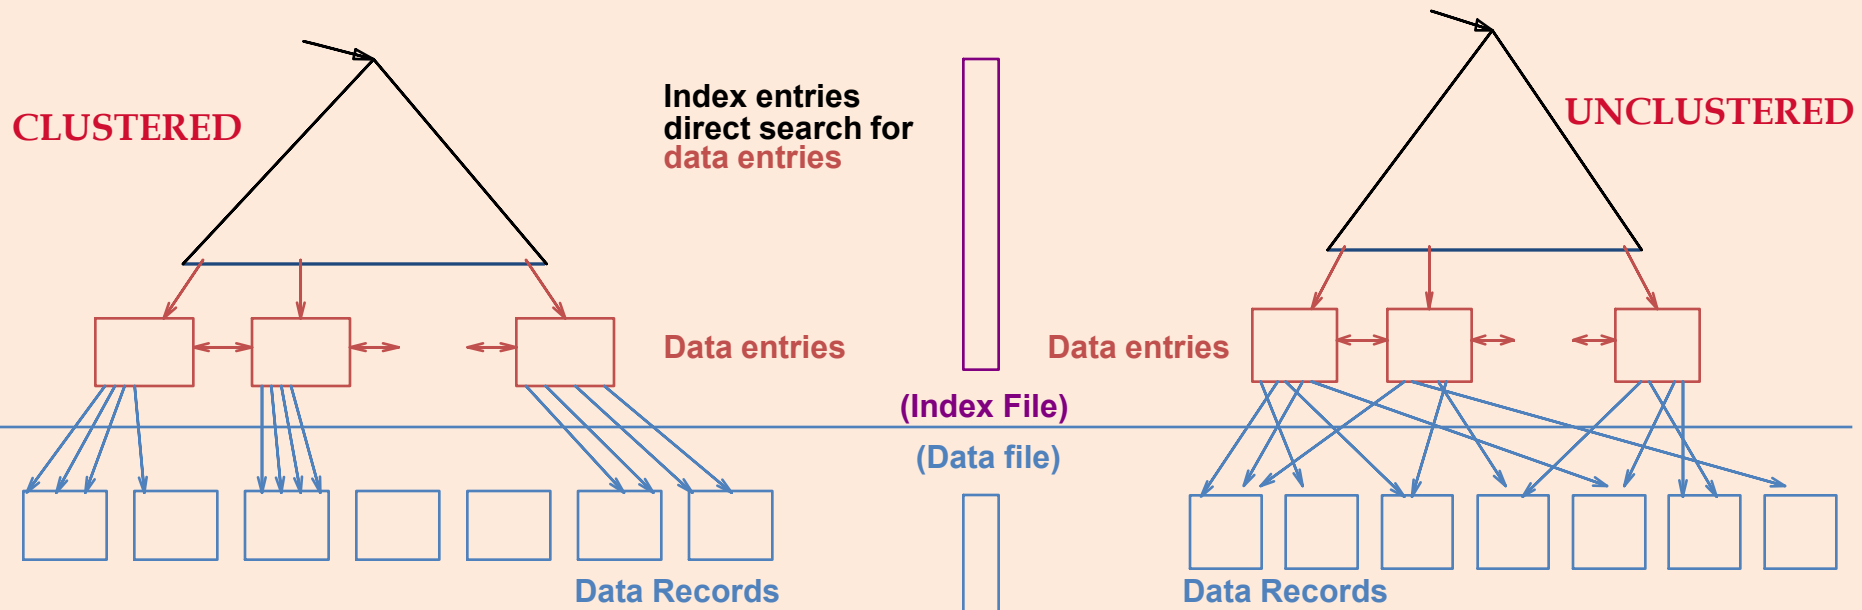
# Index Classifications

- What should be in a Data Entry  $k^*$  ?
  - Possibilities:
    - The data record itself with key value  $k$
    - $\langle k, \text{rid of data record with key value } k \rangle$
    - $\langle k, \text{list of rids of data records with key value } k \rangle$ 
      - Variable size data entries
  - Applies to any indexing technique
- Primary vs Secondary
  - **Primary index** : search key contains primary key
  - **Unique Index** : search key contains candidate key
- Clustered vs unclustered
  - **Clustered index**: order of data records same or close to order of data entries

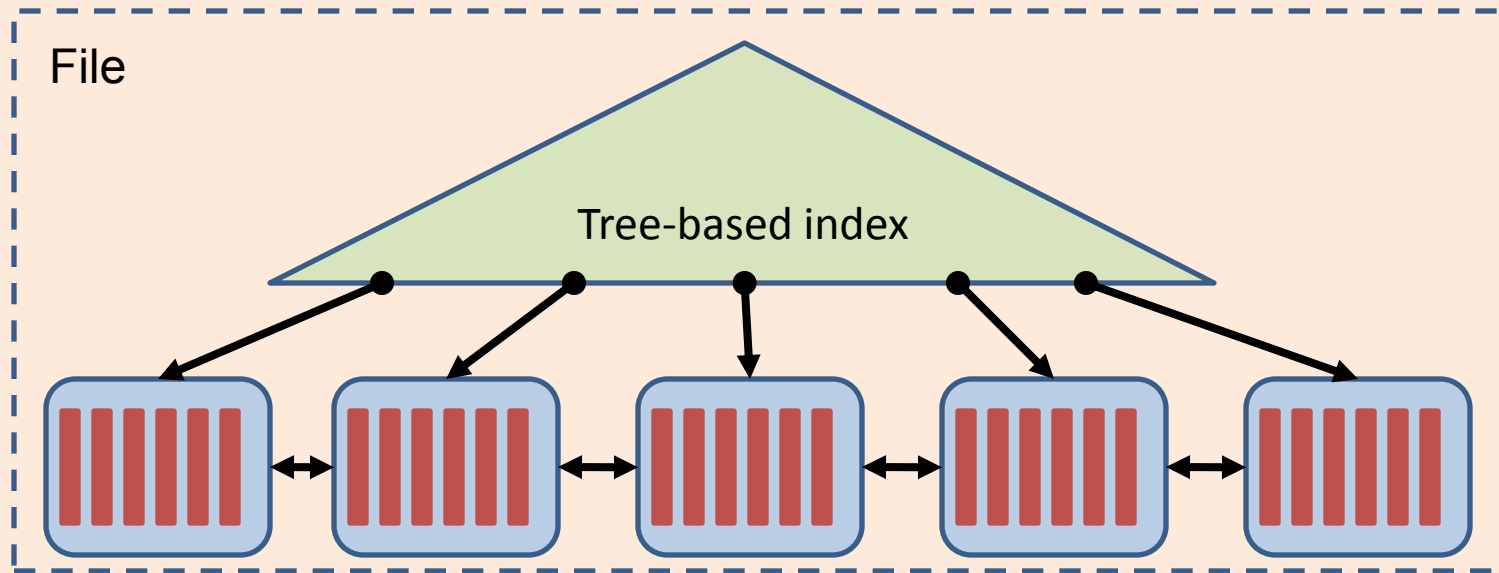


# Clustered vs Unclustered Index

- Suppose data records are stored in a Heap file.
  - To build clustered index, first sort the Heap file (with some free space on each page for future inserts).
  - Overflow pages may be needed for inserts. (Thus, order of data recs is 'close to', but not identical to, the sort order.)



# Clustered File



- An index where the data entry contains the data record itself (cf. just the key value, RID pair).
- No heap/sorted file is used, the index IS the file of record
- Steps to build a clustered file:
  - Sort data records
  - Partition into pages
  - Build the tree on the pages