

ICS 321 Data Storage & Retrieval

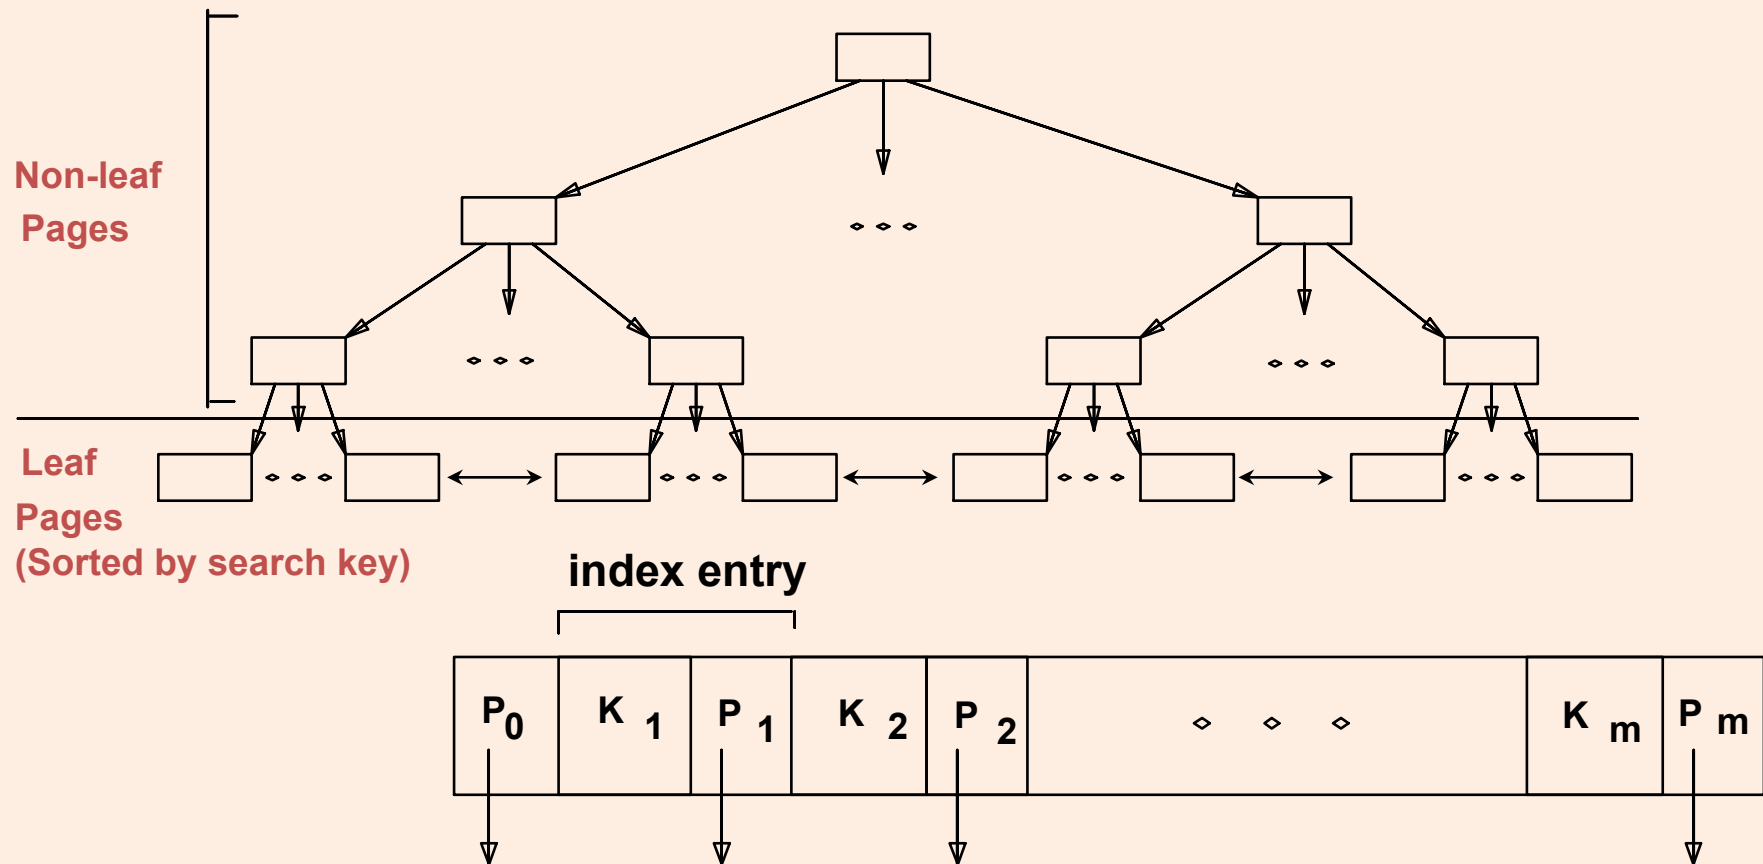
Overview of Storage & Indexing (ii)

Prof. Lipyeow Lim
Information & Computer Science Department
University of Hawaii at Manoa

Indexes

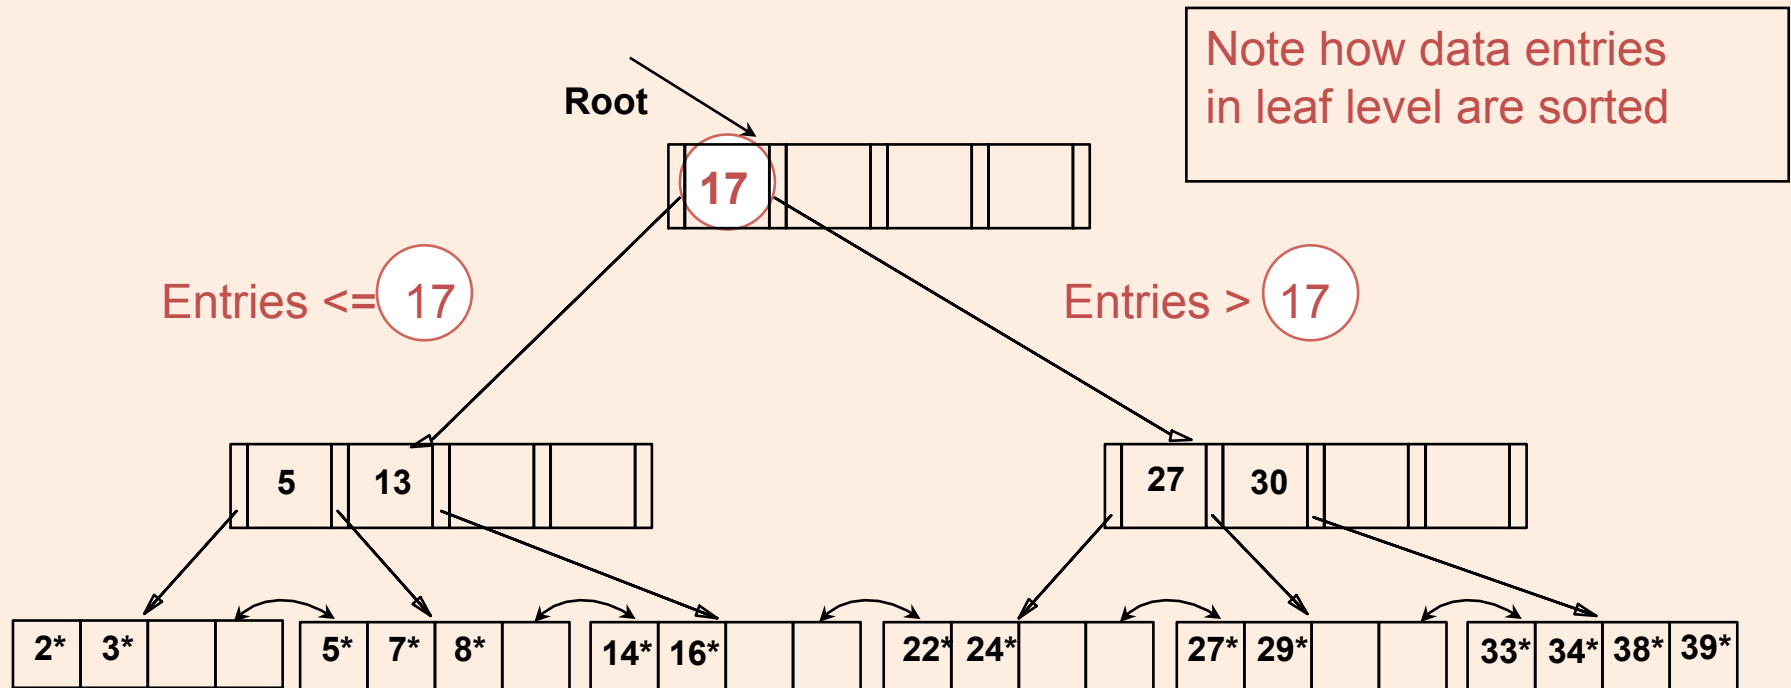
- An index on a file speeds up selections on the *search key fields* for the index.
 - Any subset of the fields of a relation can be the search key for an index on the relation.
 - *Search key* is **not** the same as *key* (minimal set of fields that uniquely identify a record in a relation).
- An index contains a collection of *data entries*, and supports efficient retrieval of all data entries k^* with a given key value k .
 - A data entry is usually in the form $\langle \text{key}, \text{rid} \rangle$
 - Given data entry k^* , we can find record with key k in at most one disk I/O. (Details soon ...)

B+ Tree Indexes



- Leaf pages contain **data entries**, and are chained (prev & next)
- A data entry typically contain a key value and a rid.
- Non-leaf pages have **index entries**; only used to direct searches:

Example B+ Tree



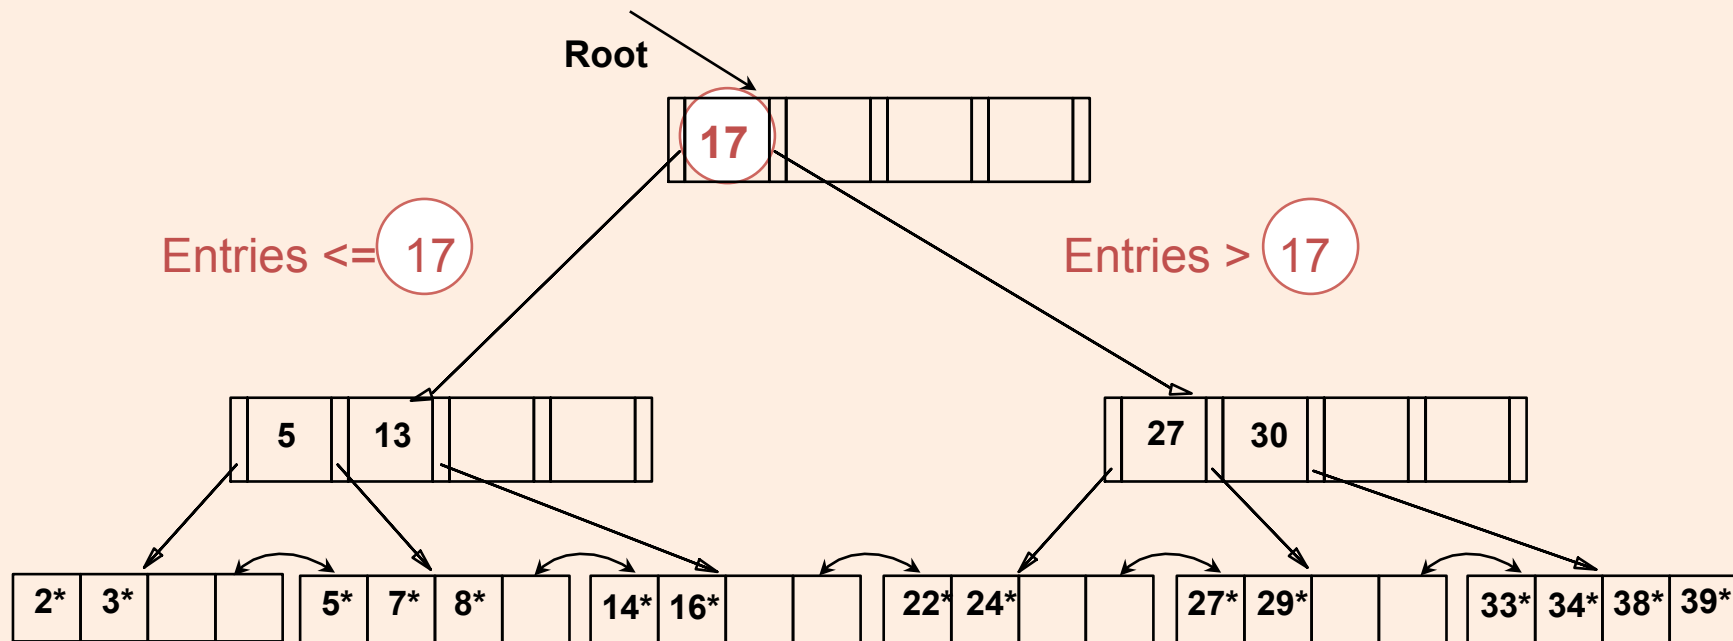
- Find 28*? 29*? All $> 15^*$ and $< 30^*$
- Insert/delete: Find data entry in leaf, then change it. Need to adjust parent sometimes.
 - And change sometimes bubbles up the tree

Point Queries using B+ Trees

SELECT *
FROM Employees
WHERE age=30

Assume heap file
data storage

- Use index to find 30*
- Request tuple from buffer manager
- If not in bufferpool, fetch page from disk

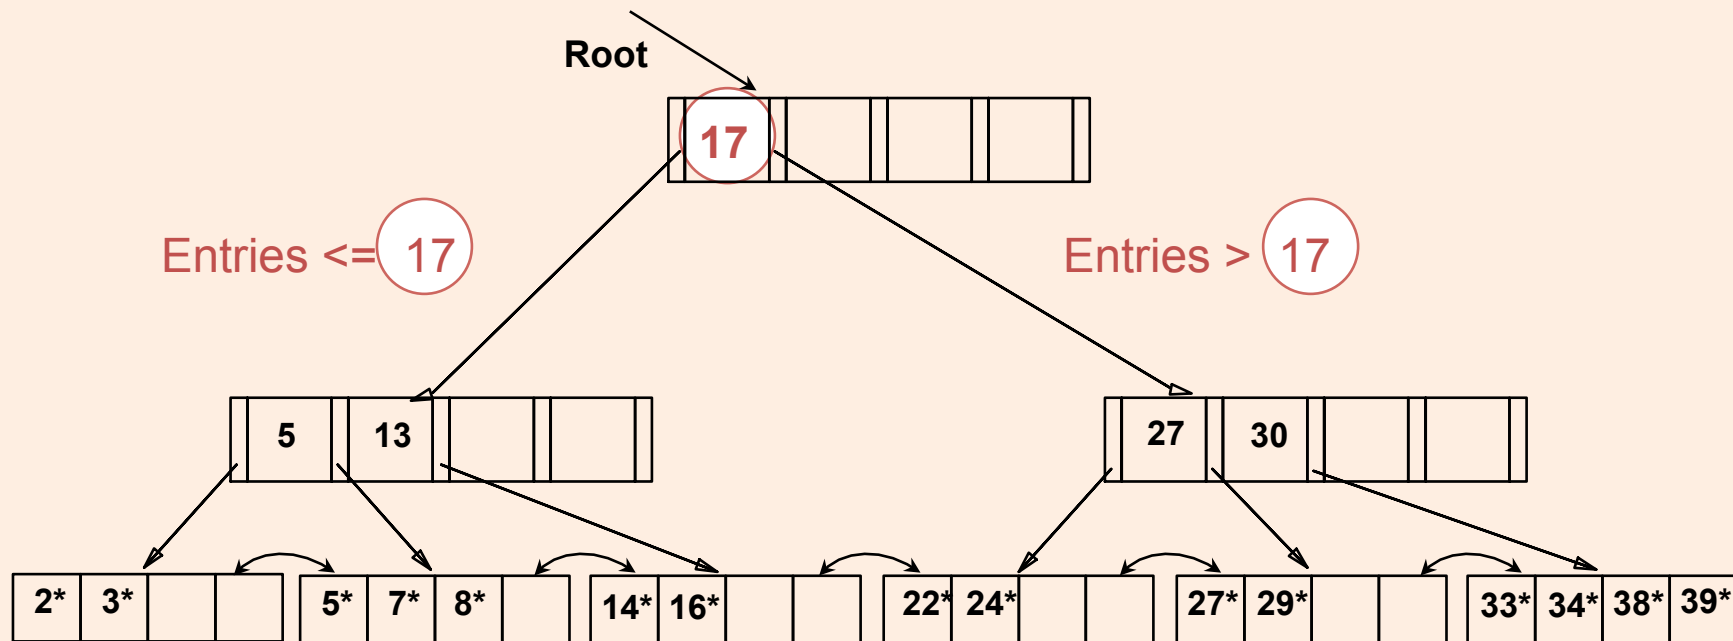


Range Queries using B+ Trees

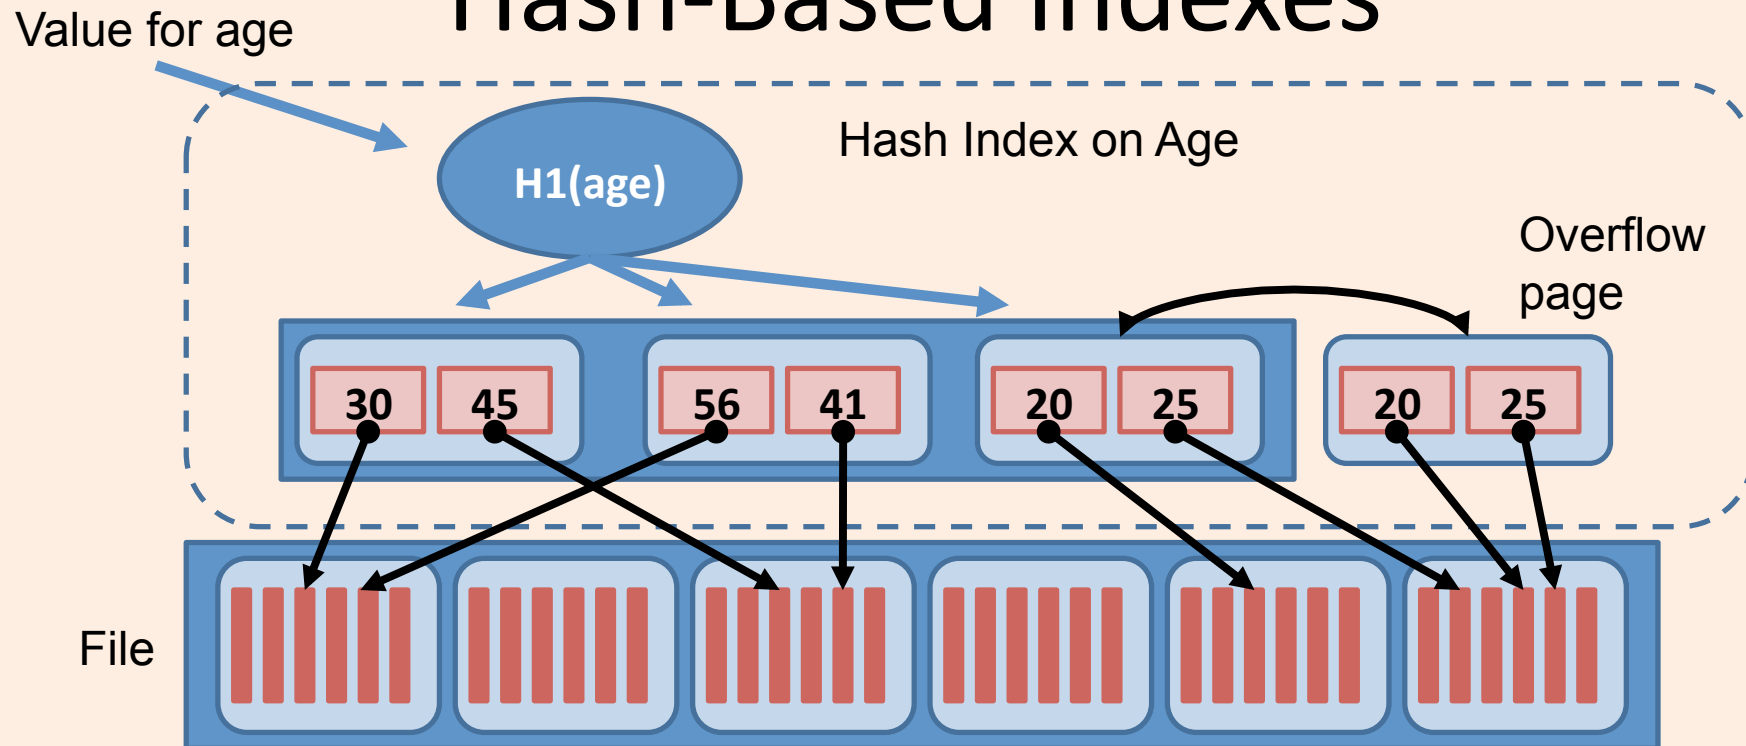
SELECT *
FROM Employees
WHERE age > 30

Assume heap file
data storage

- Use index to find 30*
- For each data entry to the right of 30*
- Request tuples from buffer manager
- If not in bufferpool, fetch page from disk



Hash-Based Indexes



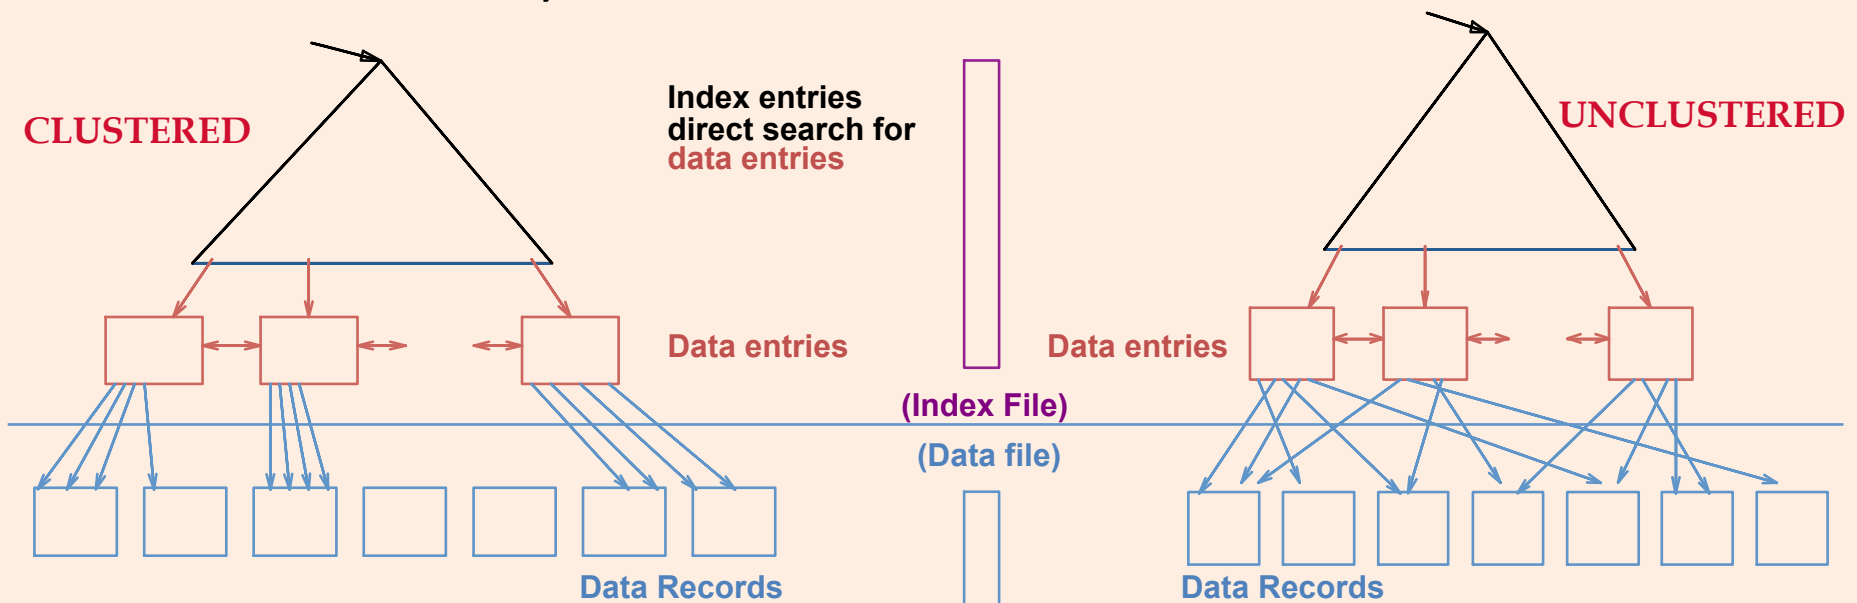
- Index is a collection of *buckets* that contain data entries
 - Bucket = *primary page* plus zero or more *overflow pages*.
- *Hashing function h*: $h(r)$ = bucket in which (data entry for) record r belongs. h looks at the *search key* fields of r .
- No “index entries” in this scheme.

Index Classifications

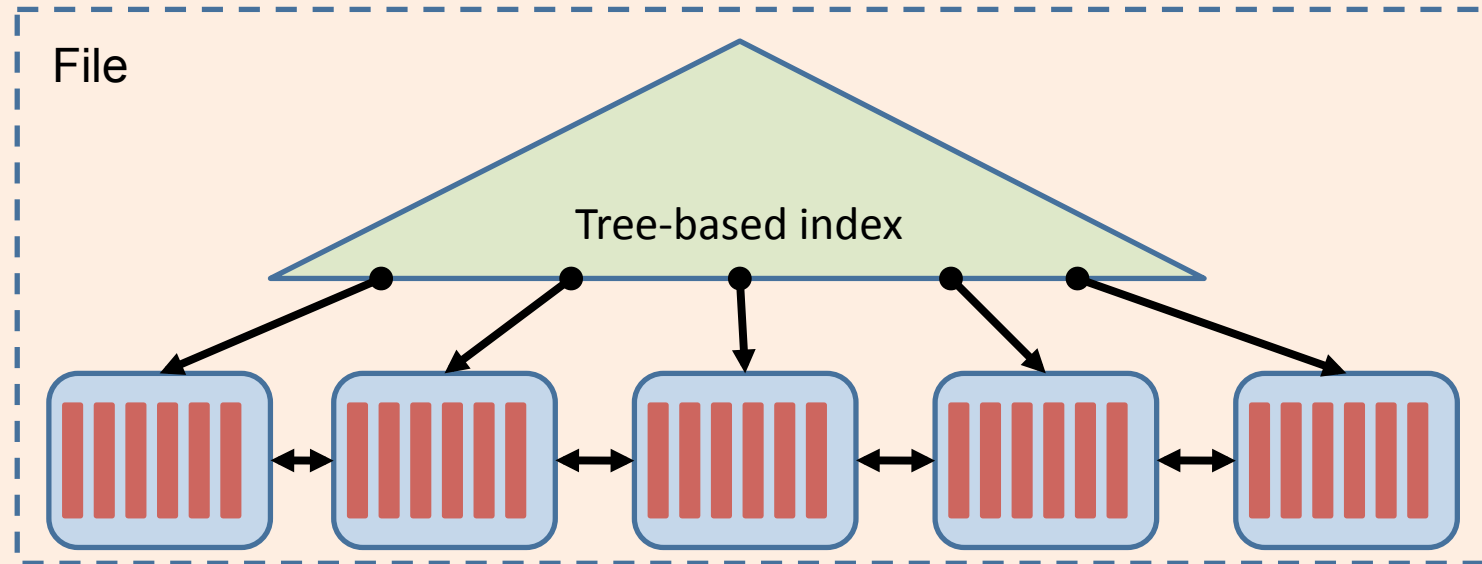
- What should be in a Data Entry k^* ?
 - Possibilities:
 - The data record itself with key value k
 - $\langle k, \text{rid of data record with key value } k \rangle$
 - $\langle k, \text{list of rids of data records with key value } k \rangle$
 - Variable size data entries
 - Applies to any indexing technique
- Primary vs Secondary
 - **Primary index** : search key contains primary key
 - **Unique Index** : search key contains candidate key
- Clustered vs unclustered
 - **Clustered index**: order of data records same or close to order of data entries

Clustered vs Unclustered Index

- Suppose data records are stored in a Heap file.
 - To build clustered index, first sort the Heap file (with some free space on each page for future inserts).
 - Overflow pages may be needed for inserts. (Thus, order of data recs is 'close to', but not identical to, the sort order.)



Clustered File



- An index where the data entry contains the data record itself (cf. just the key value, RID pair).
- No heap/sorted file is used, the index IS the file of record
- Steps to build a clustered file:
 - Sort data records
 - Partition into pages
 - Build the tree on the pages